

Lab 5: Running Hadoop Job on NYU HPC

Lab 5 outline

1. Create GitHub repo & upload dataset
2. Connect to NYU Dataproc cluster
3. Linux warm-up commands
4. Write the MapReduce job (Python)
5. Run locally to debug
6. Run on Hadoop cluster
7. Retrieve results from HDFS

Why This Lab?



HW3 Preparation

Homework 3 requires running real Hadoop jobs on this exact cluster. This lab is your practice run so nothing is new on submission day.



Hands-On Skills

You learn the complete workflow: write code → test locally → deploy on cluster → retrieve results from HDFS.



Linux commands

This mirrors how data engineers work with the linux commands

The Big Data Problem

One Laptop

Limited to 4–16 CPU cores

RAM maxes out at 8–32 GB

Disk I/O becomes bottleneck

Crashes if data is too large

Real Companies Use Clusters

Billions of logs — per second — clicks, errors, events

Billions of purchases — e-commerce transactions worldwide

Petabytes of data — 1 PB = 1,000 TB = 1,000,000 GB

Hundreds of machines — working together as ONE system

What is HPC?

HPC = High Performance Computing — many machines working as one

Discussion from lecture

Many machines, one system

Dozens to thousands of servers networked together. Your job runs across all of them simultaneously— like having thousands of CPUs.

Shared resources

CPUs, memory, and storage are pooled and shared. You request resources, the cluster allocates them, and releases them when your job finishes.

NYU Dataproc cluster

NYU provides a Google Dataproc cluster — a fully managed Hadoop/Spark environment. You connect to it via a web terminal and run jobs just like on a local machine.

What is Hadoop?

Open-source framework for storing and processing large datasets across many machines

1 HDFS

Hadoop Distributed File System

- Splits large files into 128 MB blocks
- Each block stored on multiple machines (replication = safety)
- If one machine fails, data is safe on others
- Think of it as a distributed hard drive

2 MapReduce

Distributed Computation Engine

- Breaks computation into parallel tasks
- Sends each task to the machine holding the data
- "Compute where data lives" — avoids network transfer
- Results automatically merged (reduced) at the end

What is MapReduce?

"Distributed GROUP BY" — think of it as SQL's GROUP BY running across hundreds of machines in parallel

①
MAP

Reads each row, emits a (key, value) pair

Example: Row: "Alice,2024,Milk,Grocery,3.50" → emit("Grocery", 1)

💡 Converts raw rows into structured key-value pairs Hadoop can group

②
SHUFFLE & SORT

Hadoop groups all values with the same key — you never write this code

Example: "Grocery" → [1, 1, 1, 1, 1, ...]

💡 Happens automatically between Map and Reduce — Hadoop does it behind the scenes

③
REDUCE

Receives one key + all its values, outputs the final aggregated result

Example: "Grocery", [1,1,1,...] → emit("Grocery", 40)

💡 Produces your answer — one output row per unique key

Our Dataset

shopping_data_200.csv — 200 rows of shopping transactions

Column	Type	Meaning	Example
user_id	TEXT	Unique customer identifier	e.g., U001, U042
date	DATE	Transaction date	e.g., 2024-01-15
item	TEXT	Product purchased	e.g., Milk, Laptop, Pen
category ★	TEXT	Our GROUP BY key (column index 3)	Grocery, Electronics...
price	FLOAT	Purchase amount in USD	e.g., 3.50, 299.99

Sample: U012,2024-01-15,Milk,Grocery,3.50

Goal: Count purchases per category → MapReduce GROUP BY

SQL vs MapReduce — Same Goal, Different Scale

SQL Approach

```
SELECT category, COUNT(*)  
FROM shopping_data  
GROUP BY category;
```

Easy to write & read

Great for moderate data

Runs on ONE machine only

Fails on very large datasets

MapReduce Approach

```
mapper:  row → emit(category, 1)  
shuffle: Hadoop groups all same keys  
(automatic)  
reducer: (category, [1,1,1,...]) → sum →  
output
```

Runs across many machines

Scales to petabytes

More code to write

Higher setup overhead

Step 1 – Create Your GitHub Repo

①

Create a new repository

On github.com → New → Name it exactly: lab5-hpc-mapreduce (exact name used for grading)

②

Create two folders

data/ → will contain your CSV dataset

src/ → will contain your Python MapReduce code

③

Upload the dataset

Drag data/shopping_data_200.csv into the data/ folder via GitHub web UI, then commit

④

Verify repo structure

```
lab5-hpc-mapreduce/  
    ├── data/shopping_data_200.csv  
    └── src/ (empty for now — code added later)
```

Step 2 – Connect to NYU Dataproc

Open the Dataproc Web Terminal provided by your TA, then run:

```
$ mkdir lab5
```

mkdir = make directory. Creates a new folder named "lab5" in your home directory

Keeps all your lab files organized in one place — not scattered around

```
$ cd lab5
```

cd = change directory. Moves your terminal prompt into the lab5 folder

Like double-clicking a folder on Windows/Mac. All commands now run inside this folder

```
$ git clone https://github.com/<you>/lab5-hpc-mapreduce.git
```

git clone = downloads an entire GitHub repo to the cluster machine

Copies your code and data from GitHub onto Dataproc so Hadoop can access them

```
$ cd lab5-hpc-mapreduce
```

Moves into the cloned repo folder

You must be inside the repo directory for all remaining commands to work correctly

Step 3 – Linux Warm-Up Commands

Run these to verify your files are in place and explore the dataset

```
$ ls
```

list → shows all files and folders in current directory
→ data/ src/

```
$ ls data
```

list inside data/ → shows what files are in that folder
→ shopping_data_200.csv

```
$ pwd
```

print working directory → shows your exact current path on the cluster
→ /home/user/lab5/lab5-hpc-mapreduce

```
$ head data/shopping_data_200.csv
```

show first 10 lines of the file — quick sanity check to confirm columns
→ user_id,date,item,category,price
U001,2024-01-01,Milk,Grocery,3.50 ...

```
$ wc -l data/shopping_data_200.csv
```

word count -lines → counts rows. 201 = 200 data rows + 1 header row
→ 201 data/shopping_data_200.csv

```
$ head -n 1 data/... | tr ',' '\n' | nl
```

pipe: 1st row → replace commas with newlines → number each line = find column index
→ 1 user_id 2 date 3 item 4 category 5 price

Step 3b – Search Data with grep

grep = global regular expression print — filters lines containing a pattern

```
$ grep Grocery data/shopping_data_200.csv | head
```

grep Grocery

grep = filter lines

Reads every line and keeps ONLY lines that contain the word "Grocery". All other lines are discarded. Works like Ctrl+F on the entire file.

| (pipe)

| = pipeline connector

The pipe | takes the OUTPUT of the command on the LEFT and feeds it as INPUT to the command on the RIGHT. Data flows through like water in a pipe.

head

head = show first 10

Without head, ALL matching rows would print — potentially thousands. head limits output to just 10 rows so you can quickly check the results.

Step 4 — Create the MapReduce File

```
$ nano src/mr_sales_per_category.py
```

nano is a command-line text editor built into Linux. It opens a file for editing directly inside the terminal — no GUI, no mouse needed. After pasting your code: Ctrl+X → press Y → press Enter to save and exit.

Understanding the file path:

src/

mr_

sales_per_category

.py

Folder — keeps code separate from data

"mr_" prefix = MapReduce (naming convention)

Describes what the job computes

Python file extension

After saving → commit & push to GitHub: git add src/mr_sales_per_category.py && git commit -m "add MR job" && git push

Step 4b – The MapReduce Code (Full Walkthrough)

```
from mrjob.job import MRJob
import csv

class MRSalesPerCategory(MRJob):
    def mapper(self, _, line):
        if "user_id" in line:
            return
        row = next(csv.reader([line]))
        category = row[3]
        yield category, 1

    def reducer(self, key, values):
        yield key, sum(values)

if __name__ == "__main__":
    MRSalesPerCategory.run()
```

① Import MRJob framework

mrjob = Python library that simplifies Hadoop MapReduce.
import csv = handles CSV line parsing.

② Skip header row

"user_id" in line detects the header. return exits mapper early so we don't count column names.

③ Parse CSV & emit (key, 1)

csv.reader parses the line. row[3] = 4th column = category. yield sends (category, 1) to shuffle.

④ Reducer sums all 1s

Hadoop calls reducer once per unique category with all emitted 1s.
sum(values) = total purchase count.

⑤ Entry point

.run() tells MRJob to start the MapReduce job when this script is executed directly.

Step 5 — Run Locally (Debug Mode)

Always test locally before submitting to the cluster — cluster errors are much harder to diagnose

```
$ python3 src/mr_sales_per_category.py data/shopping_data_200.csv
```

python3

Run with Python 3 interpreter (not python2 — important!)

src/mr_sales_per_category.py

Your MapReduce script relative to current directory

data/shopping_data_200.csv

Input file — no cluster flag means LOCAL mode

What happens in local mode:

- ① MRJob calls mapper() on each row of your CSV — simulating the Map phase
- ② MRJob groups and sorts output by key in memory — simulating the Shuffle phase
- ③ MRJob calls reducer() for each unique key — simulating the Reduce phase
- ④ Results print to your terminal — you see the output immediately, no HDFS needed

Step 6 – Run on Hadoop Cluster

You are now submitting to real distributed machines

```
# 1. Find the Hadoop streaming JAR
$ ls /usr/lib/hadoop-mapreduce/hadoop-streaming*.jar

# 2. Create unique output directory name (timestamp)
$ OUTDIR="lab5_out_$(date +%)"

# 3. Run the job on the cluster
$ python3 src/mr_sales_per_category.py data/shopping_data_200.csv \
    -r hadoop --hadoop-streaming-jar /usr/lib/hadoop-mapreduce/hadoop-streaming*.jar \
    --output-dir $OUTDIR --python-bin python3
```

Key flags explained:

`-r hadoop` → Tells MRJob to use real Hadoop cluster instead of local mode — the KEY switch

`--hadoop-streaming-jar` → Points to the JAR that bridges Python code and Java-based Hadoop streaming engine

`--output-dir $OUTDIR` → HDFS location for results. Uses the variable set above (unique per run)

`--python-bin python3` → Tells ALL cluster workers to use python3 (not python2) when running your script

`date +%` → Unix timestamp = seconds since 1970. Makes folder name unique every run (avoids HDFS conflict)

Step 7 – View & Retrieve Results from HDFS

HDFS is NOT your local filesystem — you need special "hadoop fs" commands to interact with it

```
$ hadoop fs -ls $OUTDIR
```

hadoop fs = Hadoop filesystem client. -ls lists files stored on HDFS in your output directory.

Output is split across multiple part-00000, part-00001... files (one per reducer task). This shows them all.

```
$ hadoop fs -getmerge $OUTDIR result.out
```

-getmerge downloads ALL part-xxxxx shards and merges them into ONE local file named result.out.

Hadoop always produces multiple output files. getmerge combines them into one for easy viewing/submission.

```
$ head result.out
```

Display first 10 lines of result.out — your final MapReduce output in (key, count) format.

Quick sanity check — verify categories and counts look correct before submitting to Canvas.

```
$ hadoop fs -rm -r $OUTDIR
```

-rm -r = remove recursively. Permanently deletes the output directory from HDFS.

CRITICAL: Hadoop CANNOT overwrite an existing output dir. You MUST delete before re-running, or use a new folder name.

Expected Output

What you should see after running head result.out

"Grocery"	40
"Electronics"	38
"Stationery"	41
"Personal Care"	39
"Clothing"	42

"Grocery" **Key = category name**

The value yielded by your mapper. mrjob wraps string keys in quotes — this is normal and expected output.

\t (tab) **Tab separator**

Hadoop uses tab (\t) between key and value, not a space or comma. If you need to parse result.out, use .split("\t").

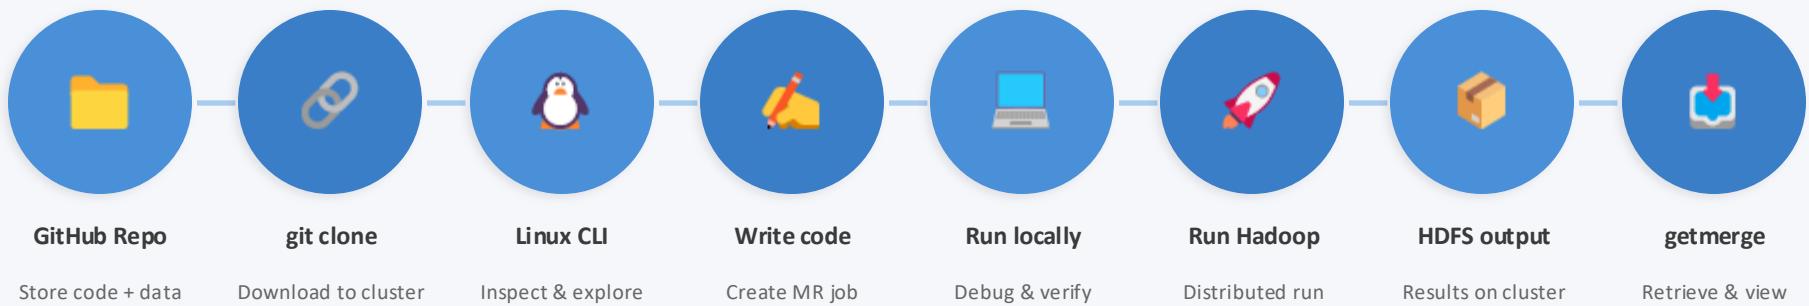
40 **Value = count**

The sum of all 1s emitted for this category. 40 means there were 40 grocery purchases in the 200-row dataset.

Order? **Output is NOT sorted**

Hadoop does not guarantee order across categories. If you need sorted output, add a sort command after getmerge.

Full Workflow – End to End



```
# Complete command sequence
git clone https://github.com/<you>/lab5-hpc-mapreduce.git
cd lab5-hpc-mapreduce
python3 src/mr_sales_per_category.py data/shopping_data_200.csv # local test
OUTDIR="lab5_out_$(date +%)"
&& python3 src/mr_sales_per_category.py data/shopping_data_200.csv -r hadoop --hadoop-streaming-jar
/usr/lib/hadoop-mapreduce/hadoop-streaming*.jar --output-dir $OUTDIR --python-bin python3
hadoop fs -getmerge $OUTDIR result.out && head result.out
```

You Are Ready for HW3!

This lab = mini version of Homework 3

✓ GitHub repos

✓ Linux CLI

✓ MapReduce code

✓ Hadoop cluster

✓ HDFS retrieval